

Utilizing Machine Learning and Spotify's API for Genre Classification in Music

Ethan David Rayala, Bodhi Harmony, Kevin Dai, Ahmadou Bamba Diouf

December 15, 2023

1 Introduction

This project's objective is to create a machine-learning model capable of identifying a song's genre by analyzing its audio features gathered through the Spotify API. This process involves collecting a range of audio features from various songs, focusing on aspects like rhythm, melody, and timbre, and then utilizing them to train a machine-learning algorithm designed to categorize songs into one of five genres. Additionally, the project assesses the model's effectiveness through an accuracy score. The outcomes of this research have potential implications for enhancing music recommendation systems, streaming services, and the broader field of music analysis.

2 Previous Work

Several past studies have applied machine learning techniques music genre classification based on sourced audio features from platforms similar to Spotify. For instance, Yang et al. in 2016 processed over 500,000 songs with the Spotify API. the team extracted key audio features and employed a wide range of machine learning methods, including logistic regression and neural networks, to determine the songs' genres. This approach, using a logistic regression model, yielded an accuracy percentage of over 70%. Similarly, Lee et al. in 2018 adopted a comparable methodology but further integrated user listening patterns and demographic data, and in turn returned an accuracy of over 80% using deep learning techniques. Kumar's 2022 project applied the K-Nearest Neighbors algorithm, achieving around 70% accuracy in music genre classification. Additionally, Ranganath's 2023 study utilized the GTZAN Genre Classification dataset, comprising 1,000 tracks across 10 genres. This project explored various audio features and employed Convolutional Neural Networks, ultimately reaching an accuracy of 92.93% in genre classification.

3 Methodology

3.1 Locations

In our Github, (<https://github.com/erayala15/COMP-562-Final-Project-Spotify>), there is a file called 'fetchspotifyapi.py' which runs the Spotify API to get the set of songs and writes them to a file called 'songdata.json'. The model is in 'model.ipynb' which reads the content of 'songdata.json'.

3.2 Intro

The primary focus of our project was the gathering and collecting of data. We relied exclusively on the Spotify API for this purpose. As detailed in the documentation available on Spotify's website, this API allows for the retrieval of various attributes of a song, based on a set of song IDs. These attributes, which range from 0.0 to 1.0, encompass several aspects of the music.

- Acousticness: A confidence measure from 0.0 to 1.0 of whether the track is acoustic.
- Danceability: Danceability describes how suitable a track is for dancing based on a combination of musical elements including tempo, rhythm stability, beat strength, and overall regularity.
- Energy: Represents a perceptual measure of intensity and activity. Typically, energetic tracks feel fast, loud, and noisy. For example, death metal has high energy, while a Bach prelude scores low on the scale.
- Instrumentalness: Predicts whether a track contains no vocals. "Ooh" and "aah" sounds are treated as instrumental in this context. Rap or spoken word tracks are clearly "vocal".
- Liveness: Detects the presence of an audience in the recording. Higher liveness values represent an increased probability that the track was performed live.
- Mode: Indicates the modality (major or minor) of a track, the type of scale from which its melodic content is derived. Major is represented by 1 and minor is 0.
- Speechiness: Speechiness detects the presence of spoken words in a track. The more exclusively speech-like the recording (e.g. talk show, audiobook, poetry), the closer to 1.0 the attribute value.
- Tempo: The overall estimated tempo of a track in beats per minute (BPM).
- Valence: Tracks with high valence sound more positive (e.g. happy, cheerful, euphoric), while tracks with low valence sound more negative (e.g. sad, depressed, angry).

Our goal is to use these specific attributes about a song to train a model to predict the genre.

3.3 Challenges

Our initial challenge involved identifying a collection of songs belonging to specific genres. Utilizing Spotify’s search API, we could specify a genre in the search parameters. However, it’s important to note that Spotify does not directly associate genres with individual songs; instead, genres are linked to artists and albums. Our approach was to use the search API to obtain song IDs and then retrospectively assign genre tags to these songs for our machine-learning model. A significant issue we faced was that the songs returned from the API searches were only loosely related to the desired genres and often did not meet the subjective standards of the genre (the ‘ear test’). Moreover, the number of songs returned for each query was inconsistent. In our preliminary tests using these songs, the model’s accuracy fluctuated between 50-70%. Notably, when we broadened our search criteria to include distinct genres like ‘classical piano’, ‘electronica’, and ‘screamo’, our model’s accuracy impressively increased to 96%.

3.4 The Playlist Solution

After attempting to locate songs via Spotify’s search API, we shifted to using Spotify’s official playlists for genre-specific song selection, still applying our method of retrospectively assigning genre tags. This strategy had two drawbacks: potential bias in the song selection and a limited sample size. Nonetheless, given our time limitations, this seemed like the best approach. The songs from these popular playlists appeared to be more representative of their respective genres, especially when subjected to the ‘ear test’, compared to those found through the search API.

3.5 Model

For our model, we chose Python and scikit-learn. Initially, we compiled the data into a data frame. This data was then divided into training and testing sets for model training and performance evaluation, respectively. We utilized scikit-learn’s `LabelBinarizer()` to transform categorical data into binary format, a requirement for most machine learning models. We then prepared a matrix with the selected attributes for training. Utilizing the `RandomForestClassifier()` function from scikit-learn, we created 100 decision trees for the training process. The model was fitted to the training data, and predictions were made on the test set.

3.6 Evaluation

To assess our model, we used the ‘score’ method from sklearn, which calculates the mean accuracy (ranging from 0.0 to 1.0) for the test data against our genre labels. We noticed that accuracy was higher with just two genres and decreased as more genres were added. By selectively including specific attributes in the model, such as danceability, speechiness, and tempo, we improved its accuracy. Interestingly, removing the ‘mode’ attribute, which represents the major or minor key of a song, slightly enhanced the accuracy.

4 Conclusion

4.1 Results

In the end, our selection included 50 songs spanning five different genres: Metal, Salsa, Jazz, Country, and Reggaeton. Impressively, our model achieved a .94 accuracy rate. This particular mix of songs revealed that our model was exceptionally proficient in identifying and distinguishing Metal and Jazz, correctly classifying every instance of these two genres. Moreover, it never mistakenly identified any song from the other genres as either Metal or Jazz.

4.2 Constraints

The primary constraint in our study was the limited size of our sample comprising of only 250 songs, with 50 reserved for testing. For future improvements, we would have either enhanced our method by either better leveraging the Spotify search API or by redesigning our approach to incorporate a broader selection of songs from Spotify’s official playlists. A promising strategy to address this issue could involve utilizing Spotify’s genre-specific and regularly updated playlists of new music over an extended time frame, considering these playlists are refreshed weekly.

4.3 Future Work / Impact

The significance of our research lies in its empirical approach to identifying the most critical attributes that define a genre based on Spotify’s metrics. Future work could expand beyond increasing the sample size to algorithmically fine-tuning the selection of genres and song attributes used in the model. Our model’s most unique aspect is its reliance on data pulled not from preexisting song databases. Future exploration could look into the importance of specific attributes in contributing to the distinguishing between genres. For instance, in differentiating classical from rap, attributes like acousticness and speechiness were more accurate than valence and liveness, suggesting the potential of certain attributes as key indicators for classification between certain genres or as universally strong predictors in general.

References

1. C. Ranganath, "Music Genre Classification using CNN," Clairvoyant Blog, 11-Apr-2019. [Online]. Available: <https://www.clairvoyant.ai/blog/music-genre-classification-using-cnn/>. [Accessed: 08-May-2023].
2. D. Lee, Y. Lee, H. Lee, and J. Lee, "Using Deep Learning and User Data to Predict Music Genres," in Proceedings of the 2018 International Conference on Information and Communication Technology Convergence (ICTC), Jeju, Korea (South), 2018, pp. 447–449.
3. scikit-learn. "RandomForestClassifier - scikit-learn 0.24.2 documentation," [Online]. Available: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>. [Accessed: 08-May-2023].
4. S. Kumar, "Music Genre Classification Project using Machine Learning Techniques," Analytics Vidhya, 29-Mar-2022. [Online]. Available: <https://www.analyticsvidhya.com/blog/2022/03/music-genre-class>